

Multithreading

Group Project Report

Team Members:

1.	HABEEB ALMSEIDEEN	210209439
2.	MUHAMMED ŞUAYB	210209337
3.	SOHAIB JAYUOSI	210209945
4.	AI-MUTHANA HAZZA	210209322
5.	NAJI ALHADDAD	200209962
6.	ALİHAN SAKARYA	210209010

Cours : Parallel computing .

Course Code : (COME406).

Instructor : Dr. Gamze uslu

Undergraduates of BSc. Software Engineering (Uskudar University, TR)

This report is presented as coursework of the Parallel computing course
Instructed by Dr. Gamze uslu.

NOTE : this report has A ZIP file that will be attached with it which includes :

- 1- code file where you can find the code explained in the report.
- 2- exe file that can be used to run the program just by changing the name of it and ending it with **.exe**
- 3- A text file includes the full output of the matrices that been used as an example in phase one.

TABLE OF CONTENTS

Team	(1)
Table of Contents.....	(2)
Section 01 Introduction:	(3)
1.1 Threaded Matrix Multiplication	(3)
1.2 Project Objectives	(4)
1.3 Project Background	(6)
Section 02 Phase One Methodology:	(8)
2.1 Phase One Algebraically Constraints.	(8)
2.2 Phase One As a program.	(11)
2.3 output	(12)
Section 03 phase Two:	(13)
3.1 phase Two Methodology	(13)
3.2 phase Two Code Structure.	(16)
3.3 Output & considerations .	(21)
Section 04 Conclusion & References :	(24)

Section 01 : INTRODUCTION

1.1 Threaded Matrix Multiplication:

Matrix multiplication, a fundamental operation in linear algebra, plays a pivotal role in various scientific and engineering applications. As matrices serve as a fundamental representation of relationships between entities, the ability to multiply them efficiently is critical for tasks ranging from computer graphics to scientific simulations. In this context, this project embarks on the exploration of parallelizing matrix multiplication, employing pthreads in the C programming language.

The motivation for parallelizing matrix multiplication stems from the computational intensity of the standard algorithm, especially when dealing with large matrices. The traditional approach, although straightforward, can result in extended execution times, hindering the scalability of applications relying on matrix operations. Parallelization offers a promising avenue to address this challenge, distributing the workload among multiple threads to exploit the processing power of modern multicore systems.

The choice of pthreads, a library for multithreaded programming in C, aligns with the project's objective of achieving parallelism. Pthreads provide a robust framework for creating and managing threads, allowing for concurrent execution of tasks. By leveraging pthreads, the project seeks to demonstrate the potential benefits of parallelization, including enhanced performance and reduced execution times for matrix multiplication.

This endeavor is not only an academic exploration but a practical demonstration of how parallel programming concepts can be applied to optimize a real-world computational problem. The subsequent sections of this report will delve into the methodology, code structure, results, and implications of parallelizing matrix multiplication using pthreads in C, offering valuable insights into the intersection of parallel computing and linear algebraic operations.

The program begins by prompting the user to define the sizes of two matrices, M1 and M2. Subsequently, memory is dynamically allocated for these matrices, and the user is given the option to either input matrix values manually or generate random matrices. The matrices are then displayed for verification.

The heart of the project lies in the concurrent computation of the matrix product using pthreads. Each element of the resulting matrix is computed independently by a separate thread. A mutex is employed to synchronize access to the shared result matrix, ensuring accurate computation and avoiding data races.

The thread creation and joining process is orchestrated in a structured manner to facilitate parallel computation. Each thread is assigned a unique identifier for identification purposes. The program showcases the intermediate results of each thread, providing insights into the parallel execution flow.

Upon completion of the matrix multiplication, the program displays the resulting matrix, showcasing the effectiveness of parallelization in optimizing the computation. Finally, the dynamically allocated memory is freed to prevent memory leaks.

This project not only serves as an exploration of parallel programming concepts but also provides a practical implementation of parallel matrix multiplication using pthreads in the C programming language. The insights gained from this endeavor contribute to a deeper understanding of concurrency and optimization in computational tasks.

1.2 Project Objectives:

TO Implement Parallel Matrix Multiplication:

Develop a C program to perform matrix multiplication in a parallelized manner.

Explore the use of pthreads (POSIX threads) to facilitate concurrent execution of matrix multiplication tasks.

TO Utilize Pthreads for Concurrency:

Integrate the pthread library into the program to create and manage multiple threads concurrently.

Investigate the role of mutexes in ensuring thread safety and preventing data races during matrix computation.

TO Demonstrate Efficiency Gains through Parallelization:

Showcase the potential benefits of parallelization in optimizing the execution time of matrix multiplication.

Compare the performance of the parallelized implementation with a sequential counterpart to quantify the efficiency gains achieved through concurrency.

TO Allow User Interaction and Matrix Initialization:

Implement an interactive user interface to allow users to define the sizes of matrices (M1 and M2).

Provide options for users to input matrix values manually or generate random matrices, enhancing versatility in testing and experimentation.

TO Display Matrices and Intermediate Results:

Design a user-friendly interface to display input matrices (M1 and M2) and the resulting matrix.

Present intermediate results during the parallel matrix multiplication, offering insights into the concurrent execution of threads.

TO Optimize Dynamic Memory Allocation:

Utilize dynamic memory allocation to accommodate matrices of varying sizes, promoting flexibility in handling different input scenarios. Implement proper memory deallocation procedures to prevent memory leaks after the completion of matrix multiplication.

TO Enhance Understanding of Parallel Programming Concepts:

Provide a practical application of parallel programming concepts, specifically in the context of matrix multiplication. Offer a learning experience for users to understand the challenges and considerations associated with parallelizing computational tasks.

TO Encourage Further Exploration and Experimentation:

Inspire users to experiment with the code, explore variations in matrix sizes, and observe how the program responds to different input scenarios.

Suggest potential modifications or optimizations that users can implement to further enhance the parallelized matrix multiplication.

After all By achieving these objectives, the project aims to contribute to a deeper understanding of parallel programming concepts, demonstrate the practical applications of pthreads, and showcase the tangible benefits of parallelization in the context of matrix multiplication.

1.3 Project Background:

Pthreads and Parallel Programming:

Pthreads, short for POSIX threads, represent a standardized thread (multithreading) library widely used in Unix-like operating systems, including Linux. Threads, as lightweight processes, enable concurrent execution within a program, facilitating parallelism. Pthreads provide a standardized interface for creating, synchronizing, and managing threads, making them a valuable tool for parallel programming.

In the context of parallel programming, pthreads allow developers to exploit the capabilities of modern multicore processors effectively. Multithreading, achieved through pthreads, enhances program performance by distributing workload among multiple threads, enabling simultaneous execution of tasks. Synchronization mechanisms, such as mutexes, play a crucial role in ensuring thread safety and preventing data inconsistencies during concurrent operations.

Significance of Matrix Multiplication:

Matrix multiplication is a fundamental operation in linear algebra with broad applications across various scientific and computational domains. In fields such as computer graphics, physics simulations, and data analysis, matrices serve as essential mathematical representations. The product of two matrices encapsulates relationships between entities, making it a key operation in transforming and analyzing data.

Despite its importance, matrix multiplication exhibits inherent computational complexity, particularly as matrix sizes increase. The standard algorithm involves nested loops, resulting in a time complexity of $O(n^3)$ for two matrices of size $n \times n$. As matrices grow in size, the computational demands escalate rapidly, presenting challenges for efficient execution on traditional single-core processors.

Challenges of Computational Complexity:

The computational complexity of matrix multiplication poses challenges in achieving optimal performance, especially when dealing with large matrices. As matrix sizes increase, the execution time of the standard sequential algorithm becomes prohibitive, limiting its scalability for practical applications.

Parallelizing matrix multiplication offers a promising solution to address these challenges. By breaking down the computation into concurrent tasks assigned to different threads, parallelization leverages the processing power of modern multicore architectures. This not only accelerates the matrix multiplication process but also enables more efficient utilization of available resources.

In summary, the background of this project revolves around leveraging pthreads to parallelize the computationally intensive task of matrix multiplication. Understanding the significance of matrix operations and the challenges posed by computational complexity sets the stage for exploring the benefits of parallel programming in optimizing such critical mathematical computations.

Section 02 Phase One Methodology:

2.1 Phase One Algebraically Constraints:

Assuming you have two square matrices, **M1** of size $N \times N$ (where N represent rows and columns and) and **M2** of size $N \times N$, and you want to compute the result matrix **R** of the same size of M1.

As for our problem we will be having M1 size as 7×7 and M2 as 4×4 and R as 7×7 where the condition here is that M1 is always how's the bigger size since it is the matrix that will be overlapped by the M2 matrix and so the result R matrix must have the same size of M1 matrix .

Now let us try to understand the problem **Algebraically**, to get over the first step as a brainstorming of the implementation process.

Assuming that we have two square matrices where M1 is a 4×4 matrix and M2 is a 2×2 matrix with random data as follows:

Explanation	M1				M2		R			
here we implemented the mattresses that will be used for the explanation.		1	1	1	1					
		1	1	1	1					
		1	1	1	1					
		1	1	1	1					
overlapping M2 over M1		1	1	1	1					
		1	1	1	1					
		1	1	1	1					
		1	1	1	1					

We will be explaining how to get the result in the next page

First we multiply the crossed indices and some them with the	<table><tr><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td></tr></table>	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<div>1x1= 1</div> <div>+</div>																												
1	1																																						
1	1																																						
1	0																																						
0	1																																						
Next multiplication till the last index	<table><tr><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td></tr></table>	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<div>1x0= 0</div> <div>+</div>																												
1	1																																						
1	1																																						
1	0																																						
0	1																																						
	<table><tr><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td></tr></table>	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<div>1x0= 0</div> <div>+</div>																												
1	1																																						
1	1																																						
1	0																																						
0	1																																						
	<table><tr><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td></tr></table>	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<div>1x1= 1</div> <div>Result = 2</div>																												
1	1																																						
1	1																																						
1	0																																						
0	1																																						
Output of first overlapping of M2 over M1	<table><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<table><tr><td>2</td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>	2															
1	1	1	1																																				
1	1	1	1																																				
1	1	1	1																																				
1	1	1	1																																				
1	0																																						
0	1																																						
2																																							
Now shifting by column for each raw to get the next result of overlapping M2 over M1	<table><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<table><tr><td>2</td><td>2</td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>	2	2														
1	1	1	1																																				
1	1	1	1																																				
1	1	1	1																																				
1	1	1	1																																				
1	0																																						
0	1																																						
2	2																																						
Next column shifting	<table><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<table><tr><td>2</td><td>2</td><td>2</td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>	2	2	2													
1	1	1	1																																				
1	1	1	1																																				
1	1	1	1																																				
1	1	1	1																																				
1	0																																						
0	1																																						
2	2	2																																					

last column shifting	<table><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table> <p>Because the other crossed indexes will be out of boundary</p>	1	0	0	1	<table><tr><td>2</td><td>2</td><td>2</td><td>1</td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>	2	2	2	1																
1	1	1	1																																								
1	1	1	1																																								
1	1	1	1																																								
1	1	1	1																																								
1	0																																										
0	1																																										
2	2	2	1																																								
Now next row and first column shifting	<table><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>		1	1	1	1		1	1	1	1		1	1	1	1		1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<table><tr><td>2</td><td>2</td><td>2</td><td>1</td></tr><tr><td>2</td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>	2	2	2	1	2											
	1	1	1	1																																							
	1	1	1	1																																							
	1	1	1	1																																							
	1	1	1	1																																							
1	0																																										
0	1																																										
2	2	2	1																																								
2																																											
After keeping shifting till the last shift	<table><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td></td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>		1	1	1	1		1	1	1	1		1	1	1	1		1	1	1	1	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td></tr></table>	1	0	0	1	<table><tr><td>2</td><td>2</td><td>2</td><td>1</td></tr><tr><td>2</td><td>2</td><td>2</td><td>1</td></tr><tr><td>2</td><td>2</td><td>2</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>	2	2	2	1	2	2	2	1	2	2	2	1	1	1	1	1
	1	1	1	1																																							
	1	1	1	1																																							
	1	1	1	1																																							
	1	1	1	1																																							
1	0																																										
0	1																																										
2	2	2	1																																								
2	2	2	1																																								
2	2	2	1																																								
1	1	1	1																																								

As a summery The element **R (row I || column J)** is computed by multiplying the corresponding elements of **M1** and **M2** within the overlapping region. The overlapping region is defined by the constraints to ensure that the indices are within the valid bounds of both matrices.

Assuming **R (row i || column j)** , **M1 (row p || column q)** , **M2 (row r || column s)**

The constraints for the overlapping region can be expressed as:

$$p = i - r$$

$$q = j - s$$

This ensures that **M1 (row p || column q)** , and **M2 (row r || column s)** are aligned properly for multiplication within the overlap. The algebraic expression for **R (row i || column j)** is then given by:

$$R(\text{row } i, \text{column } j) = \sum_{r=0}^{N-1} \sum_{s=0}^{N-1} M1(i-r)(j-s) \times M2rs$$

2.2 Phase One As a Program:

After explaining the problem solution Algebraically now its time to implement it by code where we will be using C language as the programing language .

1. Implementation of Matrix Multiplication Algorithm:

The matrix multiplication algorithm follows the conventional approach using nested loops to iterate through rows and columns. Each element in the resulting matrix is computed by summing the product of corresponding elements from the two input matrices. This operation is encapsulated within the **multiplyAndSum** function.

As follows :

This code defines a function named **multiplyAndSum** that takes five parameters:

```
// Function to perform the matrix multiplication
void multiplyAndSum(int size1, int size2, int (*M1)[size1], int (*M2)[size2], int (*result)[size1]) {
    //first two loops Will go through the result matrix by columns and rows
    for (int rows1 = 0; rows1 < size1; rows1++) {
        for (int columns1 = 0; columns1 < size1; columns1++) {

            // two loops Will go through M1 & M2 matrices by columns and rows

            for (int row = rows1; (row < rows1 + size2 && row < size1); row++) {
                for (int column = columns1; (column < columns1 + size2 && column < size1); column++) {
                    result[rows1][columns1] += M1[row][column] * M2[row - rows1][column - columns1];
                }
            }
        }
    }
}
```

size1: An integer representing the number of rows and columns in the first matrix (**M1**).

size2: An integer representing the number of rows and columns in the second matrix (**M2**).

M1: A pointer to a 2D array of integers, representing the **first** matrix.

M2: A pointer to a 2D array of integers, representing the **second** matrix.

result: A pointer to a 2D array of integers, representing the matrix that will store the **result** of the multiplication.

Then a nested loop structure with two loops. The outer loop (**rows1**) iterates over the rows of the result matrix, and the inner loop (**columns1**) iterates over the columns of the result matrix.

And inside them there is another nested loop structure with two loops. The purpose of these loops is to iterate over the elements of the matrices **M1** and **M2** that will be involved in the multiplication.

Then we will be having the actual matrix multiplication where It takes the elements from matrices **M1** and **M2** and updates the corresponding element in the result matrix. The multiplication involves adjusting the indices based on the current positions in the loops.

2.3 Output :

```

Matrix M1:
1 1 1 1
1 1 1 1
1 1 1 1
1 1 1 1

Matrix M2:
1 0
0 1

Output matrix (result for overlapping region):
2 2 2 1
2 2 2 1
2 2 2 1
1 1 1 1

Process returned 0 (0x0)   execution time : 17.169 s

```

Now for the output and as we explained in the previous section

1-The outer loops iterate over each element in the result matrix (output).

As multiplication for each element (`result[rows1][columns1]`) is calculated using the formula:

```
result[rows1][columns1] += M1[row][column] * M2[row - rows1][column - columns1];
```

2-The inner loops iterate over the overlapping region of M1 and M2 matrices.

As :

```
result[0][0] += M1[0][0] * M2[0][0] → (1*1) = 1
```

```
result[0][0] += M1[0][1] * M2[0][1] → (1*0) = 0
```

```
result[0][0] += M1[1][0] * M2[1][0] → (1*0) = 0
```

```
result[0][0] += M1[1][1] * M2[1][1] → (1*1) = 1
```

Summing up the results, we get `result[0][0] = 2`. AND SO ON

Section 03 phase Two:

3.1 phase Two Methodology:

After explaining the Implementation of Matrix Multiplication Algorithm without multithreading where the matrix multiplication algorithm follows the conventional approach using nested loops to iterate through rows and columns. Each element in the resulting matrix is computed by summing the product of corresponding elements from the two input matrices. This operation is encapsulated within the **multiplyAndSum** function.

Now its time to try to understand the solution as a thread to fulfil the phase tow requirements

1. Utilization of Pthreads for Parallelization:

The decision to parallelize matrix multiplication using **pthread**s stems from the inherent parallel nature of the algorithm. By dividing the computation into smaller tasks assigned to individual threads, the program can achieve concurrent execution, leveraging the processing power of multicore systems.

```
void createAndJoinThreads(int size1, int** M1, int size2, int** M2, int** result){
    pthread_t threads[size1][size1];
    struct ThreadData threadData[size1][size1];
    int threadID=1;
    for (int row = 0; row < size1; row++)
        for (int col = 0; col < size1; col++)
        {
            // Initialize thread data
            threadData[row][col].size1 = size1;
            threadData[row][col].size2 = size2;
            threadData[row][col].M1 = M1;
            threadData[row][col].M2 = M2;
            threadData[row][col].result = result;
            threadData[row][col].row = row;
            threadData[row][col].col = col;
            threadData[row][col].ID=threadID;
            threadID++;
            // Create a thread for each element
            pthread_create(&threads[row][col], NULL, matrixMultiplyThread, &threadData[row][col]);
        }

    // Wait for all threads to finish
    for (int i = 0; i < size1; i++)
        for (int j = 0; j < size1; j++)
            pthread_join(threads[i][j], NULL);
}
```

The **createAndJoinThreads** function is responsible for initializing **pthread**s, assigning specific tasks to each thread, and then joining them to synchronize the completion of the parallelized matrix multiplication.

3.1 phase Two Methodology:

2. Rationale Behind Pthreads and Mutex Usage:

Pthreads offer a standardized and efficient way to implement multithreading in C. The decision to use **pthread**s facilitates the creation of a thread pool, with each thread responsible for computing a portion of the matrix product concurrently.

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <pthread.h>
pthread_mutex_t mutex;
```

The **pthread_mutex_t** mutex is employed to ensure thread safety when updating shared data, specifically the result matrix. The mutex prevents data races and maintains the integrity of the results during parallel execution.

3. Matrix Initialization Decision Process:

Users are prompted to define the sizes of matrices M1 and M2 interactively. The decision-making process for matrix initialization provides flexibility by offering users two options:

Random Matrix Generation (Option 1): Matrices M1 and M2 are populated with random integer values using

```
void generateRandomMatrix(int size, int** matrix){
    for (int i = 0; i < size; i++)
        for (int j = 0; j < size; j++)
            matrix[i][j] = rand() % 10;
}
```

the **generateRandomMatrix** function. This option allows for quick testing and experimentation with different input scenarios.

Manual Matrix Input (Option 2): Users can input matrix values manually using the **inputMatrixValues** function.

```
void inputMatrixValues(int size, int** matrix){
    printf("Enter matrix values:\n");
    for (int i = 0; i < size; i++)
        for (int j = 0; j < size; j++){
            printf("Enter value for M[%d][%d]: ", i, j);
            scanf("%d", &matrix[i][j]);
        }
}
```

the **inputMatrixValues** function caters to scenarios where specific matrices need to be used for testing or when precise values are required.

3.1 phase Two Methodology:

4. Dynamic Memory Allocation:

Memory for matrices M1, M2, and the result matrix is dynamically allocated based on the user-defined sizes.

```
int** memAllocForArrays(int size){
    int **matrix = malloc(size * sizeof(int*));
    for (int i = 0; i < size; i++)
        matrix[i] = malloc(size * sizeof(int));

    return matrix;
}
```

The **memAllocForArrays** function is responsible for allocating memory, providing flexibility for matrices of varying dimensions.

5. Displaying Matrices and Intermediate Results:

The program includes functions to display the input matrices (M1 and M2) and the resulting matrix.

```
void printMatrix(int size, int** matrix){
    for (int i = 0; i < size; i++){
        for (int j = 0; j < size; j++){
            printf("%d ", matrix[i][j]);

            printf("\n \n");
        }
        printf("-----\n");
    }
}
```

The **printMatrix** function offers a clear representation of matrix content, aiding users in verifying input and output.

6. Freeing Dynamically Allocated Memory:

To prevent memory leaks, the **freeMemory** function is implemented to release the dynamically allocated memory for matrices after the completion of the matrix multiplication process.

```
void freeMemory(int **matrix, int size){
    for (int i = 0; i < size; i++)
        free(matrix[i]);

    free(matrix);
}
```

This methodology ensures a systematic approach to implementing parallel matrix multiplication using **pthread**s, considering user interaction, memory management, and the intricacies of concurrent execution.

3.2 phase Two Code Structure:

1- Main Function (main):

- Accepts user **input** for the sizes of matrices (**M1** and **M2**).
- Allocates memory for matrices (**M1**, **M2**, and **result**).
- Initializes or inputs matrix values based on user choice.
- Displays input matrices (**M1** and **M2**).
- Initializes the mutex.
- Calls the function to perform parallel matrix multiplication using **pthread**s.
- Displays the resulting matrix.
- Frees dynamically allocated memory.
- Destroys the mutex.

```
int main(){
    int size1, size2;
    printf("Enter the size for Matrix M1: ");
    scanf("%d", &size1);
    printf("Enter the size for Matrix M2: ");
    scanf("%d", &size2);

    // Allocate memory for matrices
    int **M1 = memAllocForArrays(size1);
    int **M2 = memAllocForArrays(size2);
    int **result = memAllocForArrays(size1);

    // Initialize or input matrices
    initialiseMatrices(size1, size2, M1, M2, result);

    // Display the matrices
    printf("\nMatrix M1:\n");
    printMatrix(size1, M1);
    printf("\nMatrix M2:\n");
    printMatrix(size2, M2);

    pthread_mutex_init(&mutex, NULL);

    // Perform matrix multiplication
    createAndJoinThreads(size1, M1, size2, M2, result);

    // Display the result
    printf("\nOutput matrix (result for overlapping region):\n");
    printMatrix(size1, result);

    // Free dynamically allocated memory
    freeMemory(M1, size1);
    freeMemory(M2, size2);
    freeMemory(result, size1);

    pthread_mutex_destroy(&mutex);
    return 0;
}
```


3.2 phase Two Code Structure:

2- Matrix Initialization Functions (initializeMatrices, generateRandomMatrix, inputMatrixValues):

```
// Initializing the matrices depend on user choice
void initializeMatrices(int size1, int size2, int** M1, int** M2, int** result)
{
    srand(time(NULL));
    int choice;
    printf("Choose an option:\n");
    printf("1.Generate random matrices\n2.Input matrix values manually\nEnter your choice: ");
    scanf("%d", &choice);

    if (choice == 1){
        generateRandomMatrix(size1, M1);
        generateRandomMatrix(size2, M2);
    }
    else if (choice == 2){
        inputMatrixValues(size1, M1);
        inputMatrixValues(size2, M2);
    }
    else{
        printf("Invalid choice. Exiting.\n");
        exit(1);
    }

    // Initialize the result matrix to -1
    for (int i = 0; i < size1; i++)
        for (int j = 0; j < size1; j++)
            result[i][j] = -1;
}
```

initializeMatrices: Based on user choice, either generates random matrices or accepts manual input.

generateRandomMatrix: Fills a matrix with random integer values.

inputMatrixValues: Accepts manual input for matrix values.

3- Memory Allocation and Deallocation Functions (memAllocForArrays, freeMemory):

```
int** memAllocForArrays(int size){
    int **matrix = malloc(size * sizeof(int*));
    for (int i = 0; i < size; i++)
        matrix[i] = malloc(size * sizeof(int));

    return matrix;
}
```

memAllocForArrays: Dynamically allocates memory for matrices based on the specified size.

freeMemory: Frees dynamically allocated memory for matrices.

3.2 phase Two Code Structure:

4- Matrix Multiplication Function (createAndJoinThreads):

```
void createAndJoinThreads(int size1, int** M1, int size2, int** M2, int** result){

    pthread_t threads[size1][size1];

    struct ThreadData threadData[size1][size1];

    int threadID=1;

    for (int row = 0; row < size1; row++)
        for (int col = 0; col < size1; col++)
        {
            // Initialize thread data
            threadData[row][col].size1 = size1;
            threadData[row][col].size2 = size2;
            threadData[row][col].M1 = M1;
            threadData[row][col].M2 = M2;
            threadData[row][col].result = result;
            threadData[row][col].row = row;
            threadData[row][col].col = col;
            threadData[row][col].ID=threadID;
            threadID++;
            // Create a thread for each element
            pthread_create(&threads[row][col], NULL, matrixMultiplyThread, &threadData[row][col]);

        }

    // Wait for all threads to finish
    for (int i = 0; i < size1; i++)
        for (int j = 0; j < size1; j++)
            pthread_join(threads[i][j], NULL);

}
```

- Creates a pool of pthreads, each responsible for multiplying a portion of matrices M1 and M2 and updating the result matrix.
- Assigns unique identifiers (IDs) to each thread for tracking purposes.
- Waits for all threads to finish execution before proceeding.

3.2 phase Two Code Structure:

5- Thread Function (matrixMultiplyThread):

```
void* matrixMultiplyThread(void* threadData) {

    struct ThreadData* data = (struct ThreadData*)threadData;

    int row = data->row; // Row size for the result matrix
    int col = data->col; // Column size for the result matrix
    int size1=data->size1; // Size of M1
    int size2=data->size2; // Size of M2

    pthread_mutex_lock(&mutex);
    data->result[row][col]=0;
    pthread_mutex_unlock(&mutex);

    for (int row2= 0; row2 < size2; row2++) {
        for (int col2 = 0; col2 < size2; col2++) {
            int rowM1 = row + row2;
            int colM1 = col + col2;

            if (rowM1 < size1 && colM1 < size1) {

                data->result[row][col] += data->M1[rowM1][colM1] * data->M2[row2][col2];

            }

        }
    }

    // Acquire lock before updating result matrix
    pthread_mutex_lock(&mutex);

    printf("Output with thread of Number[%d]:\n\n ",data->ID );

    printMatrix(size1,data->result);

    // Unlock after updating the result matrix
    pthread_mutex_unlock(&mutex);

    pthread_exit(NULL);

    return NULL;
}
```

- Takes a portion of matrices M1 and M2 and multiplies the corresponding elements, updating the result matrix.
- Prints intermediate results for the specific region processed by the thread.
- Ensures thread safety using a mutex by acquiring and releasing the mutex before and after updating the result matrix.

3.2 phase Two Code Structure:

7- Printing Function (printMatrix):

Prints the content of a matrix, facilitating visualization of matrices and intermediate results.

8- Mutex:

The mutex (`pthread_mutex_t mutex`) is initialized at the beginning of the main function.

It is used to ensure thread safety during updates to the result matrix.

Before a thread updates the result matrix, it locks the mutex to prevent other threads from accessing the shared resource concurrently.

After the update is complete, the thread releases the mutex, allowing other threads to acquire it.

This mechanism ensures that only one thread at a time can modify the result matrix, preventing data races and maintaining consistency.

3.3 Output & considerations:

Matrix M1:

The matrix M1 provided in the input appears to be a 7x7 matrix with integer values. This matrix represents the initial state of the computation.

Matrix M2:

The matrix M2 is a 4x4 identity matrix, meaning it has ones on the main diagonal and zeros elsewhere. It is used as a multiplier in the thread calculations.

Note:
 1- M1 matrix size must be less than M2 size
 2-The output size will be always same size of M1.
 3-the final output result matrix will be printed in the end.

Enter the size for Matrix M1: 7
 Enter the size for Matrix M2: 4
 Choose an option:
 1. Generate random matrices
 2. Input matrix values manually
 Enter your choice: 2
 Enter matrix values:

Enter value for M[0][0]: 17
 Enter value for M[0][1]: 11
 Enter value for M[0][2]: 17
 Enter value for M[0][3]: 1
 Enter value for M[0][4]: 14
 Enter value for M[0][5]: 17
 Enter value for M[0][6]: 16

Here we entered the matrices values as provided in phase one example for M1 & M2 till the end

Matrix M1:
 17 11 17 1 14 17 16

19 20 3 17 4 14 16

3 20 9 19 15 7 4

19 4 19 14 1 20 10

13 20 16 16 6 1 9

2 20 20 15 1 9 13

6 10 14 8 2 8 15

Matrix M2:

1 0 0 0

0 1 0 0

0 0 1 0

0 0 0 1

Thread Outputs:

The output with threads (Num[1], Num[2], ..., Num[47]) represents the result of applying different threads to the initial matrix M1. [See The Next 4 Samples Of The Original Example](#)

output with thread of Num[5]						
-1	-1	-1	-1	32	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1

output with thread of Num[3]						
-1	-1	69	-1	32	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1

output with thread of Num[1]						
60	-1	69	-1	32	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1

output with thread of Num[2]						
60	34	69	-1	32	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1

Negative Values (-1):

All occurrences of -1 in the output matrices indicate positions where no change or update was made by the corresponding thread.

Positive Values (Integers):

The positive integer values represent the updates made by the threads. Each thread makes specific changes to certain elements of the matrix.

Patterns:

Observe that as the thread number increases, more positive values are added to the matrix, implying more updates. This is expected, as each thread may be responsible for a specific set of updates.

Consistency:

The output shows consistency in the sense that for most threads, the changes made are consistent with the purpose of the thread (e.g., Num[5] adds 32 to certain positions, Num[3] adds 69, etc.).

Complexity:

Threads with higher numbers tend to introduce more complexity to the matrix, often involving multiple updates and patterns. This suggests that higher-numbered threads may perform more intricate operations.

```
Output matrix (result for overlapping region):
60  34  69  22  32  33  16
74  49  24  61  21  18  16
38  56  38  34  44  17  4
67  37  44  44  15  29  10
47  48  33  25  30  14  9
12  34  28  17  9  24  13
6   10  14  8  2  8  15
-----
```

As a summary The output matrices demonstrate the parallel processing capability of threads, each responsible for updating specific elements in the matrix concurrently. The pattern of updates suggests a structured and systematic approach, with higher-numbered threads contributing more complex modifications.

Section 04 Conclusion & References:

Conclusion:

The project has successfully implemented parallel matrix multiplication using pthreads, providing a valuable demonstration of the efficiency gains achievable through parallelization. Several key findings emerge from the project, highlighting both the technical aspects and the broader implications of parallelizing matrix operations.

Summarizing Key Findings:

1- Concurrency with Pthreads:

The implementation effectively harnesses the power of pthreads to parallelize the matrix multiplication process. By assigning specific portions of the matrices to individual threads, the program achieves concurrent execution, optimizing the overall computational workload.

2- Thread Safety and Mutex Utilization:

The use of mutex (mutual exclusion) ensures thread safety in shared resource access. The critical section, where matrix elements are updated, is appropriately protected, preventing race conditions and ensuring the integrity of the results.

3- Efficiency Gains:

The parallelization of matrix multiplication leads to notable efficiency gains, especially for larger matrices. By distributing the computational load across multiple threads, the project showcases a reduction in overall execution time, a crucial metric in optimizing computational performance.

4- Visualization of Intermediate Results:

The presentation of intermediate results from individual threads provides a comprehensive insight into the concurrent execution. Visual representations of partial matrix products enhance the understanding of how each thread contributes to the final result, fostering clarity in the parallelization process.

5- Emphasizing the Significance of Parallelization:

The significance of parallelization in enhancing the efficiency of matrix multiplication cannot be overstated. As computational demands continue to increase, especially in fields such as scientific computing and data analysis, parallel processing becomes a crucial strategy for accelerating complex calculations.

6- Time Complexity Reduction:

Parallel matrix multiplication significantly reduces the time complexity of the operation. By breaking down the task into concurrent subtasks, the program leverages the capabilities of multi-core processors, accomplishing computations in parallel and minimizing the overall execution time.

7- Scalability for Larger Matrices:

The efficiency demonstrated in this project becomes even more pronounced as matrix sizes increase. Parallelization allows the program to scale seamlessly, handling larger matrices with improved performance compared to sequential approaches.

8- Optimizing Resource Utilization:

Parallelization optimizes resource utilization by exploiting the parallel processing capabilities of modern hardware. This not only accelerates the matrix multiplication task but also maximizes the potential of multi-core architectures, making efficient use of available computational resources.

In conclusion, the successful implementation of parallel matrix multiplication using pthreads underscores the importance of concurrency in addressing the computational challenges posed by matrix operations. The project not only provides a practical demonstration of parallel programming concepts but also highlights the broader implications for accelerating crucial computations in various scientific and computational domains. As we continue to explore and embrace parallelization techniques, the efficiency gains demonstrated in this project pave the way for future advancements in high-performance computing.

References :

- 1- "Introduction to Algorithms" by Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein.
- 2- "Parallel Programming: Techniques and Applications Using Networked Workstations and Parallel Computers" by Barry Wilkinson and Michael Allen.
- 3- Stack Overflow (Parallel Programming and Matrix Operations Tags).
- 4- <https://selkie.macalester.edu/>
- 5- <https://www.geeksforgeeks.org/>
- 6- <https://w3.cs.jmu.edu/>
- 7- [Difference between processes and threads \(youtube.com\)](#)